

Paper Review Task 1  
Targeted Research Synthesis for Term Paper

Advances in AI-Based Toxicity Prediction for  
Drug Discovery

Submitted by,

Name: Tasnuba Tasnim

ID: 2671508

Institute of Information and Communication  
Technology

Khulna University of Engineering and  
Technology.

Submitted to,

Dr. M. A. Moyeen

Assistant Professor

Institute of Information and Communication  
Technology

Khulna University of Engineering and  
Technology.

# Table of content

---

|                                                              |          |
|--------------------------------------------------------------|----------|
| <b>I. Introduction</b>                                       | <b>2</b> |
| <b>II. State-of-the-art (SOTA)</b>                           | <b>2</b> |
| Table I: State-of-the-art comparison of twenty recent papers | 2        |
| <b>III. Conclusion</b>                                       | <b>5</b> |
| <b>IV. References</b>                                        | <b>5</b> |

# I. Introduction

Drug toxicity is a major reason why promising drug candidates fail during development. A compound may treat a disease effectively but still damage the liver, heart, kidneys, nervous system, or DNA. Traditional laboratory and animal tests remain necessary, yet they are expensive, slow, and may not fully predict human reactions. Detecting harmful compounds early is therefore important for patient safety and for reducing wasted time and cost.

Artificial intelligence (AI) supports early toxicity screening by learning patterns from chemical structures and biological data. Traditional machine-learning models use molecular descriptors and fingerprints, while newer deep-learning systems use SMILES strings, molecular graphs, attention mechanisms, and multimodal information. Recent tools predict broad ADMET properties or specific risks such as liver injury, hERG-related heart toxicity, neurotoxicity, mutagenicity, genotoxicity, and carcinogenicity. However, high accuracy on one dataset does not guarantee reliable performance on new compounds. Data imbalance, inconsistent labels, weak external validation, limited biological context, and poor interpretability remain important challenges. For practical use, AI models should provide clear explanations, uncertainty estimates, and warnings when a compound falls outside the model’s training domain.

# II. State-of-the-art (SOTA)

Table I compares twenty recent papers: four broad reviews and sixteen individual prediction platforms or model-development studies. Each work is summarized by toxicity scope, AI approach, reported result, strength, and limitation. Because the studies use different datasets and validation procedures, their scores should be interpreted within each study rather than used as a direct ranking.

Table I: State-of-the-art comparison of twenty recent papers

| Paper/Method          | Endpoint or Scope                                   | Input and AI Approach                                                                             | Reported Result | Main Strength                                                                     | Main Limitation                                                 |
|-----------------------|-----------------------------------------------------|---------------------------------------------------------------------------------------------------|-----------------|-----------------------------------------------------------------------------------|-----------------------------------------------------------------|
| Lee et al. [1]        | Review of multiple toxicity endpoints               | Reviews RF, SVM, XGBoost, neural networks, GNNs, transformers, descriptors, and molecular graphs. | No single score | Recent overview of datasets, models, endpoints, evaluation, and interpretability. | Review paper; does not experimentally validate one new model.   |
| Zhang et al. [2]      | Acute, organ-specific, genetic, and cancer toxicity | Reviews ML, deep learning, transfer learning, multimodal fusion, databases, and tools.            | No single score | Broad framework covering databases, endpoints, challenges, and future directions. | Review comparisons depend on results reported by other studies. |
| Gianquinto et al. [3] | Integrative drug safety and predictive toxicology   | Reviews QSAR, molecular modelling, ML, multi-omics,                                               | No single score | Connects computational prediction with mechanisms                                 | Review paper; no common benchmark or single model is evaluated. |

|                           |                                           |                                                                                                   |                                       |                                                                                |                                                                      |
|---------------------------|-------------------------------------------|---------------------------------------------------------------------------------------------------|---------------------------------------|--------------------------------------------------------------------------------|----------------------------------------------------------------------|
|                           |                                           | toxicokinetics, and advanced in vitro systems.                                                    |                                       | and clinical translation.                                                      |                                                                      |
| <b>Ali [4]</b>            | AI-driven drug-toxicity prediction        | Reviews QSAR, ML, deep learning, toxicity databases, algorithms, and software tools.              | No single score                       | Clear summary of approaches, challenges, regulation, and future opportunities. | Limited original quantitative comparison or external validation.     |
| <b>ADMETLab 3.0 [5]</b>   | Broad ADMET and toxicity                  | Molecular graphs with directed message-passing neural networks.                                   | hERG AUROC about 0.94                 | Predicts 119 endpoints and provides uncertainty and structural alerts.         | Benchmark results do not confirm performance in clinical settings.   |
| <b>ProTox 3.0 [6]</b>     | Organ, clinical, and general toxicity     | ML, chemical similarity, molecular targets, initiating events, and adverse outcome pathways.      | Endpoint-dependent                    | Broad coverage with some mechanistic information.                              | Dataset size, quality, and performance vary across endpoints.        |
| <b>Deep-PK [7]</b>        | 35 toxicity and pharmacokinetic endpoints | Graph neural network using molecular structures and physicochemical descriptors.                  | Endpoint-dependent                    | Highlights molecular subgraphs that influence predictions.                     | Limited dose, exposure, patient, and clinical information.           |
| <b>InterDILI [8]</b>      | Drug-induced liver injury                 | Fingerprints and descriptors with RF, LightGBM, logistic regression, and neural networks.         | AUROC 0.97; AUPRC 0.95; accuracy 0.90 | Strong performance with feature and attention explanations.                    | Performance decreased under a different data split.                  |
| <b>DILI Predictor [9]</b> | Drug-induced liver injury                 | Fingerprints, descriptors, predicted pharmacokinetic properties, proxy labels, and random forest. | AUROC 0.63                            | Combines predicted and experimental information in an accessible tool.         | Moderate discrimination and a narrow endpoint.                       |
| <b>GeoDILI [10]</b>       | Drug-induced liver injury                 | Three-dimensional molecular graphs with a geometry-based graph neural network.                    | AUROC 0.908; accuracy 0.975           | Atom-level explanations of possible toxic molecular regions.                   | Needs stronger independent, temporal, and prospective validation.    |
| <b>OvA-QSTR [11]</b>      | Drug-induced liver injury                 | PaDEL descriptors with Bayesian networks, decision trees, and random forests.                     | AUPRC 0.718-0.869                     | Relatively simple and understandable machine-learning models.                  | Hand-crafted descriptors may miss complex mechanisms.                |
| <b>Rao et al. [12]</b>    | Severity of drug-induced liver injury     | Physicochemical descriptors and off-target profiles with RF, SVM, neural networks,                | AUROC 0.88; specificity 0.90          | Predicts severity and includes off-target information.                         | Depends on the availability and accuracy of curated off-target data. |

|                             |                                                 |                                                                                                   |                                     |                                                              |                                                            |
|-----------------------------|-------------------------------------------------|---------------------------------------------------------------------------------------------------|-------------------------------------|--------------------------------------------------------------|------------------------------------------------------------|
|                             |                                                 | and logistic regression.                                                                          |                                     |                                                              |                                                            |
| <b>hERGBost [13]</b>        | hERG inhibition and heart toxicity              | Molecular descriptors and fingerprints with XGBoost classification and regression.                | Accuracy 0.814; MCC 0.614; R2 0.622 | Predicts both toxicity class and continuous IC50 values.     | Limited explanation of why a molecule is risky.            |
| <b>hERGAT [14]</b>          | hERG inhibition                                 | Molecular graphs, fingerprints, descriptors, graph attention networks, and gated recurrent units. | AUROC 0.907; AUPRC 0.904            | Atom-level and molecule-level attention explanations.        | hERG inhibition is only one mechanism of cardiotoxicity.   |
| <b>AttenhERG [15]</b>       | hERG inhibition                                 | Molecular graph representation with an AttentiveFP graph neural network.                          | AUROC 0.835; MCC 0.543              | Provides uncertainty and atom-level attention.               | Lower discrimination than several other hERG models.       |
| <b>DMFGAM [16]</b>          | hERG inhibition                                 | Extended-connectivity and atom-pair fingerprints combined with graph attention networks.          | AUROC 0.894; accuracy 0.817         | Combines molecular-graph and fingerprint information.        | Does not include clinical, metabolic, dose, or omics data. |
| <b>NeuTox 2.0 [17]</b>      | BBB effects, neuronal damage, and neurotoxicity | Molecular graphs, fingerprints, descriptors, DMPNN, MFBERT, and multimodal feature fusion.        | AUROC up to 0.971                   | Predicts several related neurotoxicity outcomes.             | Most input features are still structure-derived.           |
| <b>muTOX-AL [18]</b>        | Ames mutagenicity                               | ECFP fingerprints, MACCS keys, descriptors, 2D-CNN, and active learning.                          | AUROC 0.909                         | Strong performance with fewer labelled training samples.     | Requires repeated expert or laboratory annotation.         |
| <b>Fournier et al. [19]</b> | Genotoxicity across several assays              | 2D descriptors and MACCS keys with similarity search and extremely randomized trees.              | AUROC about 0.921-0.972             | Covers several genotoxicity assays.                          | No complete executable prediction system was reported.     |
| <b>DCAMCP [20]</b>          | Carcinogenicity                                 | Multiple molecular fingerprints with capsule networks and self-attention routing.                 | AUROC 0.793; accuracy 0.718         | Uses attention and reports balanced classification measures. | Moderate performance and limited carcinogenicity data.     |

*Note: The first four rows are review papers; the remaining sixteen rows describe individual platforms or model-development studies. AUROC and AUPRC values closer to 1 indicate stronger discrimination. Accuracy is the proportion of correct predictions, and MCC ranges from -1 to 1. Results from different datasets, endpoints, and validation procedures are not directly comparable*

### III. Conclusion

AI-based toxicity prediction has progressed from descriptor-based machine learning to graph neural networks, attention, active learning, and multimodal systems. These tools can identify possible safety risks earlier, but no single model is reliable for every endpoint. Current limitations include inconsistent datasets, weak generalization, limited biological information, and unclear model decisions. Future systems should combine molecular structure with dose, metabolism, off-target effects, and biological data while providing uncertainty and applicability warnings. AI should support, not replace, laboratory testing and clinical judgement. With transparent validation and human oversight, it can help reduce late-stage failures and support safer drug development.

### IV. References

- [1] H. Lee, J. Kim, J.-W. Kim, and Y. Lee, "Recent advances in AI-based toxicity prediction for drug discovery," *Frontiers in Chemistry*, vol. 13, Art. no. 1632046, 2025, doi: 10.3389/fchem.2025.1632046.
- [2] R. Zhang, H. Wen, Z. Lin, B. Li, and X. Zhou, "Artificial intelligence-driven drug toxicity prediction: Advances, challenges, and future directions," *Toxics*, vol. 13, no. 7, Art. no. 525, 2025, doi: 10.3390/toxics13070525.
- [3] E. Gianquinto et al., "Toward integrative predictive toxicology: Advanced methods for drug toxicity and safety prediction," *WIREs Computational Molecular Science*, vol. 16, Art. no. e70065, 2026, doi: 10.1002/wcms.70065.
- [4] S. M. Ali, "Artificial intelligence driven approaches to predicting drug toxicity: Challenges and future direction," *Egyptian Society of Clinical Toxicology Journal*, vol. 13, no. 2, pp. 1-16, 2025.
- [5] L. Fu et al., "ADMETlab 3.0: An updated comprehensive online ADMET prediction platform enhanced with broader coverage, improved performance, API functionality and decision support," *Nucleic Acids Research*, vol. 52, no. W1, pp. W422-W431, 2024, doi: 10.1093/nar/gkae236.
- [6] P. Banerjee, E. Kemmler, M. Dunkel, and R. Preissner, "ProTox 3.0: A webserver for the prediction of toxicity of chemicals," *Nucleic Acids Research*, vol. 52, no. W1, pp. W513-W520, 2024, doi: 10.1093/nar/gkae303.
- [7] Y. Myung, A. G. C. de Sa, and D. B. Ascher, "Deep-PK: Deep learning for small molecule pharmacokinetic and toxicity prediction," *Nucleic Acids Research*, vol. 52, no. W1, pp. W469-W475, 2024, doi: 10.1093/nar/gkae254.
- [8] S. Lee and S. Yoo, "InterDILI: Interpretable prediction of drug-induced liver injury through permutation feature importance and attention mechanism," *Journal of Cheminformatics*, vol. 16, Art. no. 1, 2024, doi: 10.1186/s13321-023-00796-8.
- [9] S. Seal et al., "Improved detection of drug-induced liver injury by integrating predicted and experimental data," *Chemical Research in Toxicology*, vol. 37, no. 8, pp. 1290-1305, 2024, doi: 10.1021/acs.chemrestox.4c00015.
- [10] W. X. Wu et al., "GeoDILI: A robust and interpretable model for drug-induced liver injury prediction using graph neural network-based molecular geometric representation," *Chemical Research in Toxicology*, vol. 36, no. 11, pp. 1717-1730, 2023, doi: 10.1021/acs.chemrestox.3c00199.

- [11] F. K. Celik and G. Karaduman, "Machine learning-based prediction of drug-induced hepatotoxicity: An OvA-QSTR approach," *Journal of Chemical Information and Modeling*, vol. 63, no. 15, pp. 4602-4614, 2023, doi: 10.1021/acs.jcim.3c00687.
- [12] M. Rao et al., "AI/ML models to predict the severity of drug-induced liver injury for small molecules," *Chemical Research in Toxicology*, vol. 36, no. 7, pp. 1129-1139, 2023, doi: 10.1021/acs.chemrestox.3c00098.
- [13] M.-S. Yu et al., "hERGBost: A gradient boosting model for quantitative IC50 prediction of hERG channel blockers," *Computers in Biology and Medicine*, vol. 184, Art. no. 109416, 2025, doi: 10.1016/j.combiomed.2024.109416.
- [14] D. Lee and S. Yoo, "hERGAT: Predicting hERG blockers using graph attention mechanism through atom- and molecule-level interaction analyses," *Journal of Cheminformatics*, vol. 17, Art. no. 11, 2025, doi: 10.1186/s13321-025-00957-x.
- [15] T. B. Yang et al., "AttenhERG: A reliable and interpretable graph neural network framework for predicting hERG channel blockers," *Journal of Cheminformatics*, vol. 16, Art. no. 143, 2024, doi: 10.1186/s13321-024-00940-y.
- [16] Y. M. Wang et al., "Identification of vital chemical information via visualization of graph neural networks," *Briefings in Bioinformatics*, vol. 24, no. 1, Art. no. bbac577, 2023, doi: 10.1093/bib/bbac577.
- [17] X. D. Pang et al., "NeuTox 2.0: A hybrid deep learning architecture for screening potential neurotoxicity of chemicals based on multimodal feature fusion," *Environment International*, vol. 195, Art. no. 109244, 2025, doi: 10.1016/j.envint.2024.109244.
- [18] H. Y. Xu et al., "Deep active learning with high structural discriminability for molecular mutagenicity prediction," *Communications Biology*, vol. 7, Art. no. 1071, 2024, doi: 10.1038/s42003-024-06758-6.
- [19] M. Fournier et al., "In silico genotoxicity prediction by similarity search and machine learning algorithm: Optimization and validation of the method for high energetic materials," *Propellants, Explosives, Pyrotechnics*, vol. 48, no. 4, Art. no. e202200259, 2023, doi: 10.1002/prop.202200259.
- [20] Z. Chen et al., "DCAMCP: A deep learning model based on capsule network and attention mechanism for molecular carcinogenicity prediction," *Journal of Cellular and Molecular Medicine*, vol. 27, no. 20, pp. 3117-3126, 2023, doi: 10.1111/jcmm.17889.